

Project 4: California Housing Price Prediction

MainCode

Step1: Import all libraries

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

Step2: Load the data

Step2.1: Read the “housing.csv” file from the folder into the program

```
housingData = pd.read_csv('housing.csv')
```

Step2.2: Print first few rows of this data

```
print('Print first few rows of this data - ')
```

```
print()
```

```
print(housingData.head())
```

Step2.3: Extract input (X) and output (y) data from the dataset

```
X = housingData.iloc[:, :-1].values
```

```
y = housingData.iloc[:, [-1]].values
```

Step3: Handle missing values:

Fill the missing values with the mean of the respective column

```
from sklearn.preprocessing import Imputer  
missingValueImputer = Imputer()  
X[:, :-1] = missingValueImputer.fit_transform(X[:, :-1])  
y = missingValueImputer.fit_transform(y)
```

Step4: Encode categorical data:

Convert categorical column in the dataset to numerical data

```
from sklearn.preprocessing import LabelEncoder  
X_labelencoder = LabelEncoder()  
X[:, -1] = X_labelencoder.fit_transform(X[:, -1])
```

Step5: Split the dataset: Split the data into

80% training dataset and 20% test dataset

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y,  
                                                    test_size = 0.2,  
                                                    random_state = 0)
```

Step6: Standardize data: Standardize training and test datasets

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
y_train = scaler.fit_transform(y_train)
y_test = scaler.transform(y_test)
```

```
#####
####
```

```
""Task1: Perform Linear Regression""
```

```
#####
####
```

```
# Task1.1: Perform Linear Regression on training data
```

```
from sklearn.linear_model import LinearRegression
linearRegression = LinearRegression()
linearRegression.fit(X_train, y_train)
```

```
# Task1.2: Predict output for test dataset using the fitted model
```

```
predictionLinear = linearRegression.predict(X_test)
```

```
# Task1.3: Print root mean squared error (RMSE) from Linear Regression
```

```
from sklearn.metrics import mean_squared_error
mseLinear = mean_squared_error(y_test, predictionLinear)
print('Root mean squared error (RMSE) from Linear Regression = ')
```

```
print(mseLinear)
```

```
#####  
####
```

```
"""Task2: Perform Decision Tree Regression"""
```

```
#####  
####
```

```
# Task2.1: Perform Decision Tree Regression on training data
```

```
from sklearn.tree import DecisionTreeRegressor
```

```
DTregressor = DecisionTreeRegressor()
```

```
DTregressor.fit(X_train, y_train)
```

```
# Task2.2: Predict output for test dataset using the fitted model
```

```
predictionDT = DTregressor.predict(X_test)
```

```
# Task2.3: Print root mean squared error from Decision Tree Regression
```

```
from sklearn.metrics import mean_squared_error
```

```
mseDT = mean_squared_error(y_test, predictionDT)
```

```
print('Root mean squared error from Decision Tree Regression = ')
```

```
print(mseDT)
```

```
#####  
####
```

```
'''Task3: Perform Random Forest Regression'''
```

```
#####  
####
```

```
# Task3.1: Perform Random Forest Regression on training data
```

```
from sklearn.ensemble import RandomForestRegressor
```

```
RFRegressor = RandomForestRegressor()
```

```
RFRegressor.fit(X_train, y_train)
```

```
# Task3.2: Predict output for test dataset using the fitted model
```

```
predictionRF = RFRegressor.predict(X_test)
```

```
# Task3.3: Print root mean squared error from Random Forest Regression
```

```
from sklearn.metrics import mean_squared_error
```

```
mseRF = mean_squared_error(y_test, predictionRF)
```

```
print('Root mean squared error from Random Forest Regression = ')
```

```
print(mseRF)
```

```
#####  
####
```

```
'''Task4: Bonus exercise:
```

```
    Perform Linear Regression with one independent variable'''
```

```
#####  
####
```

```
# Task4.1: Extract just the median_income column from the
# independent variables (from X_train and X_test)
```

```
X_train_median_income = X_train[:, [7]]
X_test_median_income = X_test[:, [7]]
```

```
# Task4.2: Perform Linear Regression to predict housing values
# based on median_income
```

```
from sklearn.linear_model import LinearRegression
linearRegression2 = LinearRegression()
linearRegression2.fit(X_train_median_income, y_train)
```

```
# Task4.3: Predict output for test dataset using the fitted model
```

```
predictionLinear2 = linearRegression2.predict(X_test_median_income)
```

```
# Task4.4: Plot the fitted model for training data as well as
# for test data to check if the fitted model satisfies the test data
```

```
# Task4.4.1: let us visualize the Training set
```

```
plt.scatter(X_train_median_income, y_train, color = 'green')
plt.plot (X_train_median_income,
          linearRegression2.predict(X_train_median_income), color = 'red')
```

```
plt.title ('compare Training result - median_income / median_house_value')
plt.xlabel('median_income')
plt.ylabel('median_house_value')
plt.show()
```

Task4.4.2: let us visualize the Testing set

```
plt.scatter(X_test_median_income, y_test, color = 'blue')
plt.plot (X_train_median_income,
          linearRegression2.predict(X_train_median_income), color = 'red')
plt.title ('compare Testing result - median_income / median_house_value')
plt.xlabel('median_income')
plt.ylabel('median_house_value')
plt.show()
```

```
#####
####
'''          End          '''
#####
####
```

Code Snippet followed a Screenshot of the output

Q1. Load the data and Print first few rows of this data

Step2: Load the data

Step2.1: Read the “housing.csv” file from the folder into the program

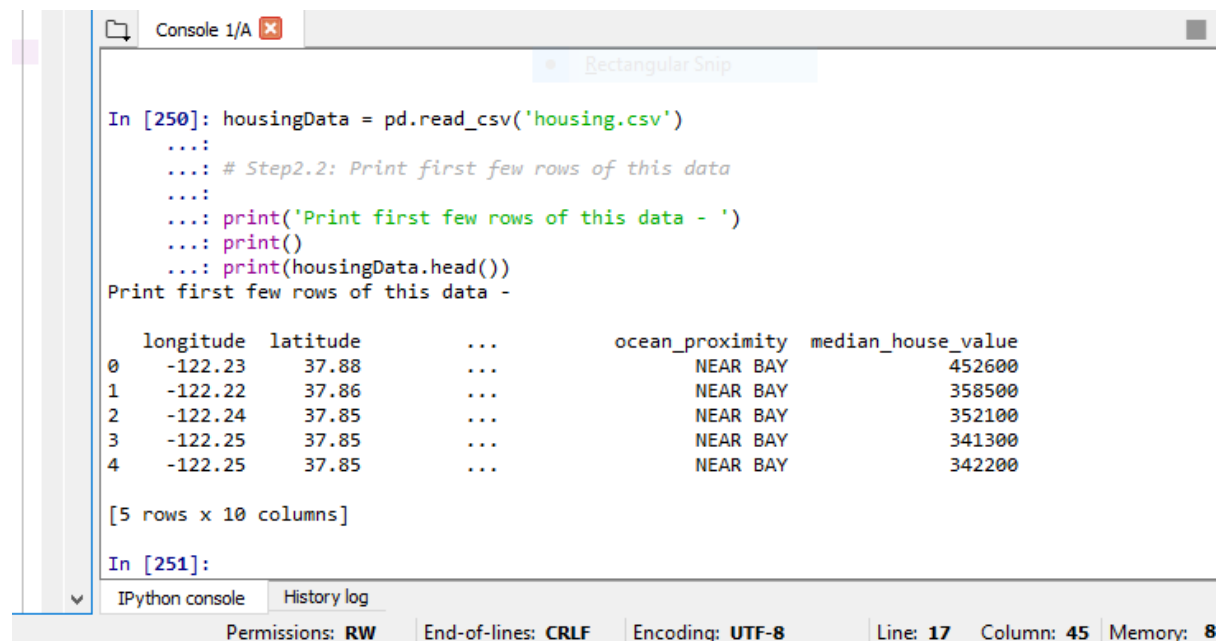
```
housingData = pd.read_csv('housing.csv')
```

Step2.2: Print first few rows of this data

```
print('Print first few rows of this data - ')
```

```
print()
```

```
print(housingData.head())
```



The screenshot shows an IPython console window with the following content:

```
In [250]: housingData = pd.read_csv('housing.csv')
...:
...: # Step2.2: Print first few rows of this data
...:
...: print('Print first few rows of this data - ')
...: print()
...: print(housingData.head())
Print first few rows of this data -
```

	longitude	latitude	...	ocean_proximity	median_house_value
0	-122.23	37.88	...	NEAR BAY	452600
1	-122.22	37.86	...	NEAR BAY	358500
2	-122.24	37.85	...	NEAR BAY	352100
3	-122.25	37.85	...	NEAR BAY	341300
4	-122.25	37.85	...	NEAR BAY	342200

```
[5 rows x 10 columns]

In [251]:
```

At the bottom of the console, there are tabs for 'IPython console' and 'History log'. Below the tabs, the status bar shows: 'Permissions: RW', 'End-of-lines: CRLF', 'Encoding: UTF-8', 'Line: 17', 'Column: 45', and 'Memory: 8'.

Q2. Separate features and labels, deal with missing value, Encode categorical data, Split the dataset into training and testing set, standardized the data

Step2.3: Extract input (X) and output (y) data from the dataset

```
X = housingData.iloc[:, :-1].values
```

```
y = housingData.iloc[:, [-1]].values
```

Step3: Handle missing values:

Fill the missing values with the mean of the respective column

```
from sklearn.preprocessing import Imputer
```

```
missingValueImputer = Imputer()
```

```
X[:, :-1] = missingValueImputer.fit_transform(X[:, :-1])
```

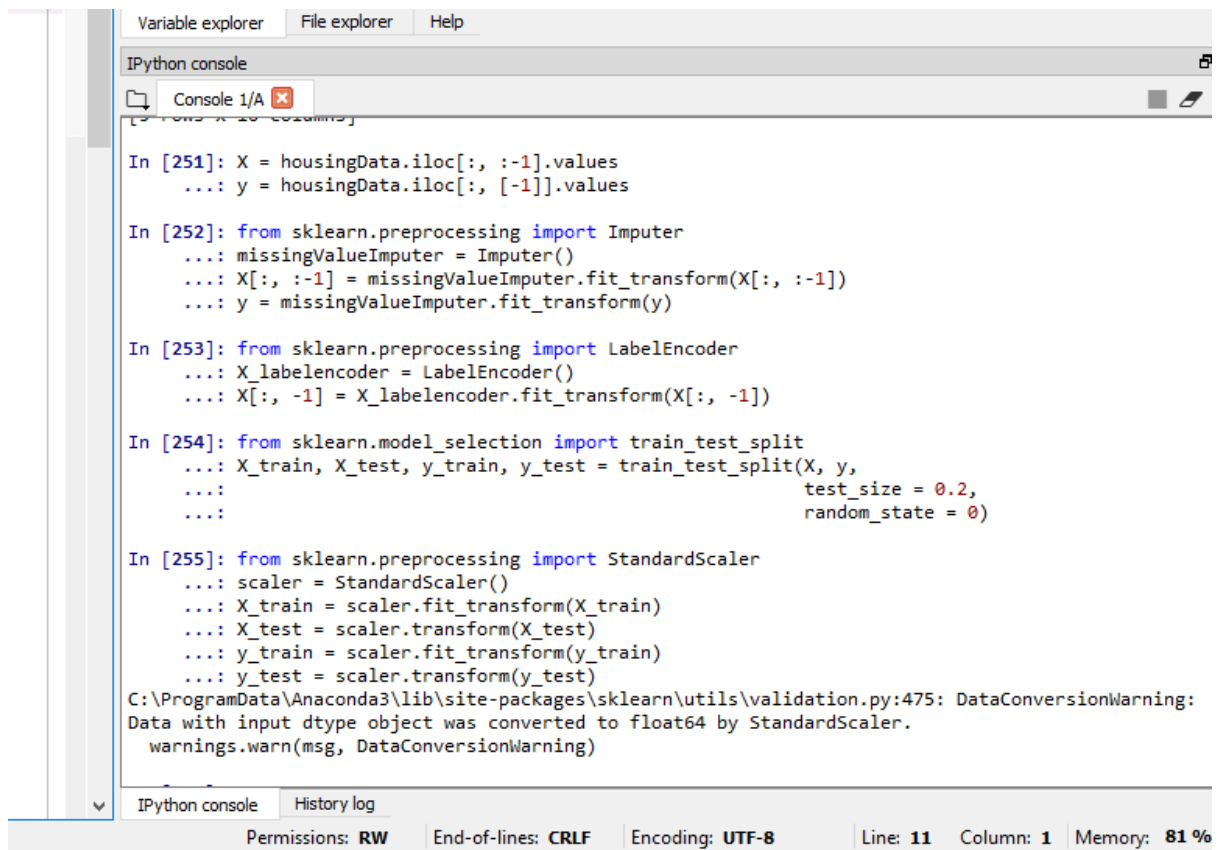
```
y = missingValueImputer.fit_transform(y)
```

Step4: Encode categorical data:


```
# Convert categorical column in the dataset to numerical data
from sklearn.preprocessing import LabelEncoder
X_labelencoder = LabelEncoder()
X[:, -1] = X_labelencoder.fit_transform(X[:, -1])

# Step5: Split the dataset: Split the data into
# 80% training dataset and 20% test dataset
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size = 0.2,
                                                    random_state = 0)

# Step6: Standardize data: Standardize training and test datasets
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
y_train = scaler.fit_transform(y_train)
y_test = scaler.transform(y_test)
```



```
Variable explorer  File explorer  Help

IPython console
Console 1/A [X]

In [251]: X = housingData.iloc[:, :-1].values
...: y = housingData.iloc[:, [-1]].values

In [252]: from sklearn.preprocessing import Imputer
...: missingValueImputer = Imputer()
...: X[:, :-1] = missingValueImputer.fit_transform(X[:, :-1])
...: y = missingValueImputer.fit_transform(y)

In [253]: from sklearn.preprocessing import LabelEncoder
...: X_labelencoder = LabelEncoder()
...: X[:, -1] = X_labelencoder.fit_transform(X[:, -1])

In [254]: from sklearn.model_selection import train_test_split
...: X_train, X_test, y_train, y_test = train_test_split(X, y,
...:                                                    test_size = 0.2,
...:                                                    random_state = 0)

In [255]: from sklearn.preprocessing import StandardScaler
...: scaler = StandardScaler()
...: X_train = scaler.fit_transform(X_train)
...: X_test = scaler.transform(X_test)
...: y_train = scaler.fit_transform(y_train)
...: y_test = scaler.transform(y_test)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:475: DataConversionWarning:
Data with input dtype object was converted to float64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)

IPython console  History log
Permissions: RW  End-of-lines: CRLF  Encoding: UTF-8  Line: 11  Column: 1  Memory: 81 %
```

Q3. Task1: Perform Linear Regression

Task1.1: Perform Linear Regression on training data

```
from sklearn.linear_model import LinearRegression
```

```
linearRegression = LinearRegression()
```

```
linearRegression.fit(X_train, y_train)
```

Task1.2: Predict output for test dataset using the fitted model

```
predictionLinear = linearRegression.predict(X_test)
```

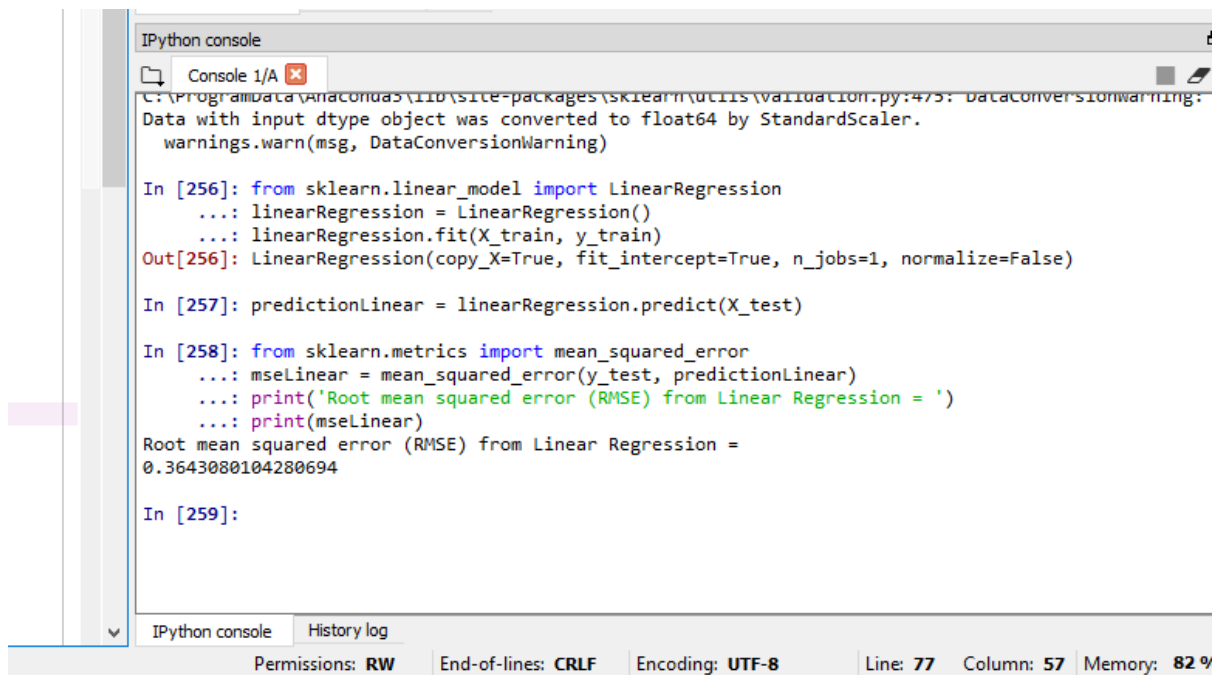
Task1.3: Print root mean squared error (RMSE) from Linear Regression

```
from sklearn.metrics import mean_squared_error
```

```
mseLinear = mean_squared_error(y_test, predictionLinear)
```

```
print('Root mean squared error (RMSE) from Linear Regression = ')
```

```
print(mseLinear)
```



```
IPython console
C:\ProgramData\Anaconda5\lib\site-packages\sklearn\utils\validation.py:475: DataConversionWarning:
Data with input dtype object was converted to float64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)

In [256]: from sklearn.linear_model import LinearRegression
...: linearRegression = LinearRegression()
...: linearRegression.fit(X_train, y_train)
Out[256]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)

In [257]: predictionLinear = linearRegression.predict(X_test)

In [258]: from sklearn.metrics import mean_squared_error
...: mseLinear = mean_squared_error(y_test, predictionLinear)
...: print('Root mean squared error (RMSE) from Linear Regression = ')
...: print(mseLinear)
Root mean squared error (RMSE) from Linear Regression =
0.3643080104280694

In [259]:
```

IPython console History log

Permissions: RW End-of-lines: CRLF Encoding: UTF-8 Line: 77 Column: 57 Memory: 82 %

Q4. Task2: Perform Decision Tree Regression

Task2.1: Perform Decision Tree Regression on training data

```
from sklearn.tree import DecisionTreeRegressor
```

```
DTRegressor = DecisionTreeRegressor()
```

```
DTRegressor.fit(X_train, y_train)
```

Task2.2: Predict output for test dataset using the fitted model

```
predictionDT = DTRegressor.predict(X_test)
```

Task2.3: Print root mean squared error from Decision Tree Regression

```
from sklearn.metrics import mean_squared_error
```

```
mseDT = mean_squared_error(y_test, predictionDT)
```

```
print('Root mean squared error from Decision Tree Regression = ')
```

```
print(mseDT)
```

```
...: print('Root mean squared error (RMSE) from Linear Regression = ')
...: print(mseLinear)
Root mean squared error (RMSE) from Linear Regression =
0.36430880104280694

In [259]: from sklearn.tree import DecisionTreeRegressor
...: DTRegressor = DecisionTreeRegressor()
...: DTRegressor.fit(X_train, y_train)
Out[259]:
DecisionTreeRegressor(criterion='mse', max_depth=None, max_features=None,
                      max_leaf_nodes=None, min_impurity_decrease=0.0,
                      min_impurity_split=None, min_samples_leaf=1,
                      min_samples_split=2, min_weight_fraction_leaf=0.0,
                      presort=False, random_state=None, splitter='best')

In [260]: predictionDT = DTRegressor.predict(X_test)

In [261]: from sklearn.metrics import mean_squared_error
...: mseDT = mean_squared_error(y_test, predictionDT)
...: print('Root mean squared error from Decision Tree Regression = ')
...: print(mseDT)
Root mean squared error from Decision Tree Regression =
0.337687895444884

In [262]:
```

Q5. Task3: Perform Random Forest Regression

Task3.1: Perform Random Forest Regression on training data

```
from sklearn.ensemble import RandomForestRegressor
```

```
RFRRegressor = RandomForestRegressor()
```

```
RFRRegressor.fit(X_train, y_train)
```

Task3.2: Predict output for test dataset using the fitted model

```
predictionRF = RFRRegressor.predict(X_test)
```

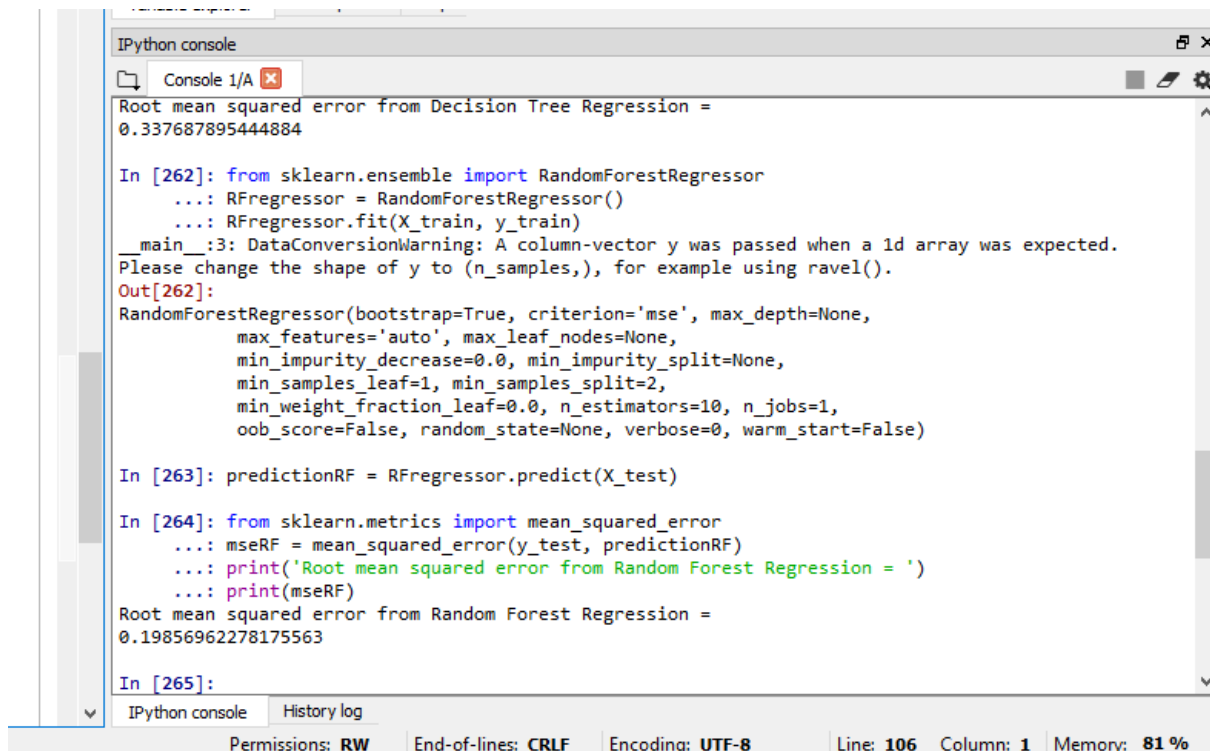
Task3.3: Print root mean squared error from Random Forest Regression

```
from sklearn.metrics import mean_squared_error
```

```
mseRF = mean_squared_error(y_test, predictionRF)
```

```
print('Root mean squared error from Random Forest Regression = ')
```

```
print(mseRF)
```



```
IPython console
Console 1/A
Root mean squared error from Decision Tree Regression =
0.337687895444884

In [262]: from sklearn.ensemble import RandomForestRegressor
...: RFRegressor = RandomForestRegressor()
...: RFRegressor.fit(X_train, y_train)
__main__:3: DataConversionWarning: A column-vector y was passed when a 1d array was expected.
Please change the shape of y to (n_samples,), for example using ravel().
Out[262]:
RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
oob_score=False, random_state=None, verbose=0, warm_start=False)

In [263]: predictionRF = RFRegressor.predict(X_test)

In [264]: from sklearn.metrics import mean_squared_error
...: mseRF = mean_squared_error(y_test, predictionRF)
...: print('Root mean squared error from Random Forest Regression = ')
...: print(mseRF)
Root mean squared error from Random Forest Regression =
0.19856962278175563

In [265]:
```

IPython console History log

Permissions: RW End-of-lines: CRLF Encoding: UTF-8 Line: 106 Column: 1 Memory: 81 %

Q6. Task4: Bonus exercise: Perform Linear Regression with one independent variable

Task4.1: Extract just the median_income column from the

independent variables (from X_train and X_test)

```
X_train_median_income = X_train[:, [7]]
```

```
X_test_median_income = X_test[:, [7]]
```

Task4.2: Perform Linear Regression to predict housing values

based on median_income

```
from sklearn.linear_model import LinearRegression
```

```
linearRegression2 = LinearRegression()
```

```
linearRegression2.fit(X_train_median_income, y_train)
```

Task4.3: Predict output for test dataset using the fitted model

```
predictionLinear2 = linearRegression2.predict(X_test_median_income)
```

```

...: print(mseRF)
Root mean squared error from Random Forest Regression =
0.19856962278175563

In [265]: X_train_median_income = X_train[:, [7]]
...: X_test_median_income = X_test[:, [7]]

In [266]: from sklearn.linear_model import LinearRegression
...: linearRegression2 = LinearRegression()
...: linearRegression2.fit(X_train_median_income, y_train)
Out[266]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)

In [267]: predictionLinear2 = linearRegression2.predict(X_test_median_income)

In [268]: plt.scatter(X_train_median_income, y_train, color = 'green')

```

Task4.4: Plot the fitted model for training data as well as

for test data to check if the fitted model satisfies the test data

Task4.4.1: **let us visualize the Training set**

```
plt.scatter(X_train_median_income, y_train, color = 'green')
```

```
plt.plot (X_train_median_income,
```

```
        linearRegression2.predict(X_train_median_income), color = 'red')
```

```
plt.title ('compare Training result - median_income / median_house_value')
```

```
plt.xlabel('median_income')
```

```
plt.ylabel('median_house_value')
```

```
plt.show()
```



Task4.4.2: let us visualize the Testing set

```
plt.scatter(X_test_median_income, y_test, color = 'blue')
```

```
plt.plot (X_train_median_income,
```


```
         linearRegression2.predict(X_train_median_income), color = 'red')
```

```
plt.title ('compare Testing result - median_income / median_house_value')
```

```
plt.xlabel('median_income')
```

```
plt.ylabel('median_house_value')
```

```
plt.show()
```



****End ****
